

Script 1 — The Pitch

~65 seconds · projected deck (FacilityTrustDesk-PitchDeck.pdf) · blue tag = advance to that slide

0–8s Intro SLIDE 1

SAY "Hello Hack-Nation! We're Mossab and Léo, and for Challenge four — Data Legend — we built Facility Trust Desk: the trust layer for Indian healthcare."

8–20s The problem SLIDE 2

SAY "In India, a hospital's ICU is often a claim, not a capability. Families drive six hours to a hospital that promised intensive care — and find a locked door. Planners don't lack data: they have ten thousand facility records. What they lack is evidence they can act on."

20–36s What we built SLIDE 3

SAY "We built Facility Trust Desk: the trust layer for that data. Every capability is treated as a claim to verify. Our engine scores each claim on independent corroboration — structured equipment and procedures against self-reported narrative — filters out 'proposed' and 'under construction', and attaches the exact sentences behind every verdict."

36–50s Honesty as a feature SLIDE 4

SAY "And the app doubts itself, on purpose. An independent validator audits our own scores — it overturned 378 of our own 'corroborated' ratings. Uncertainty bands tell planners what's solid versus speculative. Empty regions are labeled 'unknown', never 'empty' — a data desert is not a medical desert."

50–65s The close SLIDE 5

SAY "Humans stay in charge: every override is signed and remembered. Built entirely on Databricks Free Edition — Apps, serverless SQL, foundation-model embeddings, MLflow tracing, Genie. Ten thousand messy records, turned into decisions a planner can defend. That's how data becomes care."

Script 2 — The Live Demo

60 seconds · app on screen · yellow boxes = what you click

BEFORE YOU START — 2-MINUTE CHECKLIST

- Wake the app 20 min early (Free Edition sleeps) and load every tab once.
- Signed in as admin · sidebar scenario: "Maternity push — Rajasthan Q3".
- "Find facilities" open · capability **ICU** · state **Rajasthan** pre-selected.
- 2–3 facilities already shortlisted (so the Decision Brief is rich).
- Browser zoom ~110% · other tabs closed · notifications muted.

0–10s Flash recap + hook

DO Screen already shows the ICU · Rajasthan list with the stat tiles. Don't touch anything yet.

SAY "Remember — in this data, an ICU is a claim, not a verified fact. This is the app that tells the difference. Here: every ICU claim in Rajasthan, ranked by evidence. Green: independently corroborated. Amber: claimed only. Gray: we honestly don't know."

10–21s The receipts

DO Click a green **CORROBORATED** facility → expand "Evidence, gaps & review". Pick one **WITH** a red validator banner.

SAY "Every verdict shows its receipts — the exact sentences from the record, field by field, plus what we still don't know. And look: our own validator disagrees with this rating — the app audits itself and says so, out loud."

20–32s Humans take over

DO Same card: override form → new status + note "Field visit May 2026 — ICU confirmed" → Save override. Keep talking during the reload.

SAY "A field worker who knows better overrides the machine — signed, timestamped, stored for the whole team. The header counts every time a human corrected this app. We wear it as a badge."

32–44s Data desert ≠ medical desert

DO Switch to the "Medical deserts" tab → show the 4 tiles → scroll to the map.

SAY "At district level, we join the official NFHS-5 health survey. Solid red: real unmet need, proven by external indicators. Hollow gray: our records are just empty there. Two different colors, because 'no data' is not 'no hospitals' — and confusing them sends help to the wrong place."

44–54s The decision you take home

DO Switch to "Shortlist & decisions" → click "Download decision brief (.md)" on the prepared scenario.

SAY "When the planner is done, one click exports a Decision Brief: verdicts, verbatim evidence, uncertainty bands, the validator's objections, and every human override — a decision the team can defend."

54–60s The technical close

DO Point at "How was this computed?" in the sidebar — no need to click.

SAY "All of it live on Databricks Free Edition — and the full reasoning is one click away, traced end-to-end in MLflow. Thank you."

Q&A Moves & Key Numbers

If the jury digs deeper — the answer, and what to click

"Show me free-text search."

SAY "Semantic search over all ten thousand profiles, beyond our 8 fixed capabilities. Matches are leads, not verdicts."

DO *Top search bar → type "burns unit" → results show trust chips per facility.*

"Why no trained model?"

SAY "No ground truth exists to train against. A transparent rule engine is auditable sentence by sentence — and our override log is exactly the labeled data a future model needs."

"How do you know your scorer is right?"

SAY "Ten hand-validated cases in the repo, plus an independent validator that overturned 378 of our own ratings. We publish our residual limitations too."

"What's Databricks-native here?"

SAY "Apps, serverless SQL, Delta Sharing, foundation-model embeddings via ai_query, Genie, MLflow tracing, and a provisioned Lakebase instance for OLTP persistence."

"Can I ask the data myself?"

SAY "Yes — in plain English."

DO *Sidebar → "Ask the data in plain English" (Genie) → ask "How many corroborated ICU facilities in Maharashtra?" → answers ~152, SQL visible.*

NUMBER	WHAT IT IS
10,088	facilities in the dataset
29,047	claims (facility × capability) assessed
7,505	corroborated by 2+ independent sources
378	ratings our own validator overturned
1,360	facilities whose GPS contradicts their PIN code
755	districts classified (81% joined to NFHS-5)
35	normalized states (from 254 dirty values)